

Manticore

CS769 Natural Language Processing

Topics

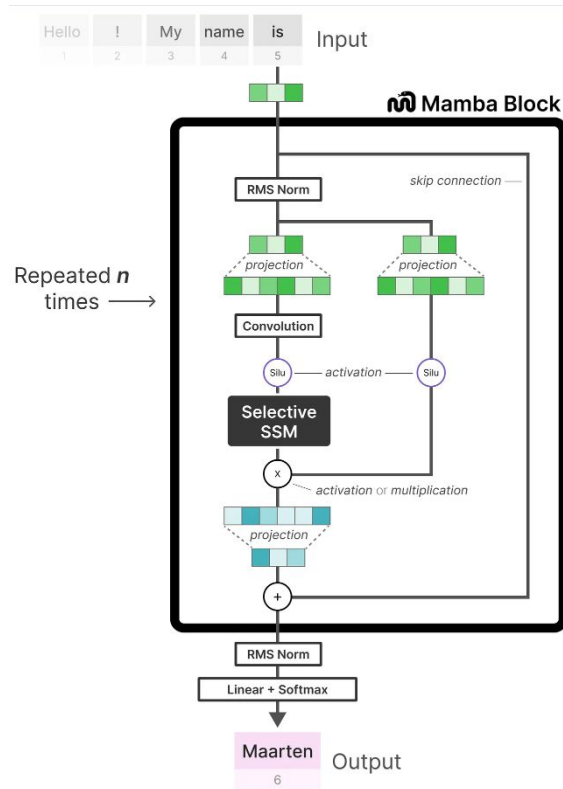
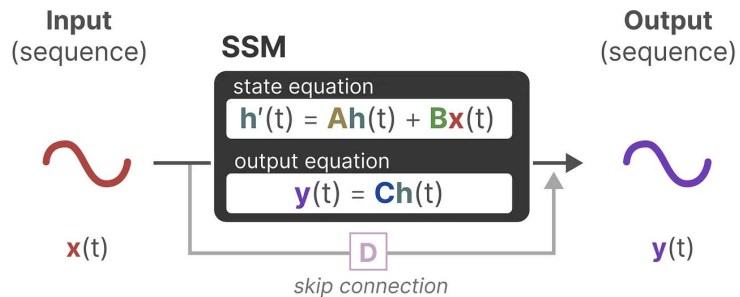

- Motivation
- Manticore framework
- Implementation
 - Results on IMDB
 - Results on MAD tasks

Motivation


- Transformer models have dominated the NLP arena
 - But computational cost is quadratic in sequence length (attention layer)
 - **Transformer dominance:** Challenged by new subquadratic architectures (e.g. State Space Models, Mamba)
- State Space models
 - Mamba
- Hybrid architectures: Combine strengths of diverse models but suffer from:
 - **Manual design:** Inefficient and heuristic-driven
 - **Pre-training challenges:** Limited ability to reuse pretrained components (Neural architecture search)



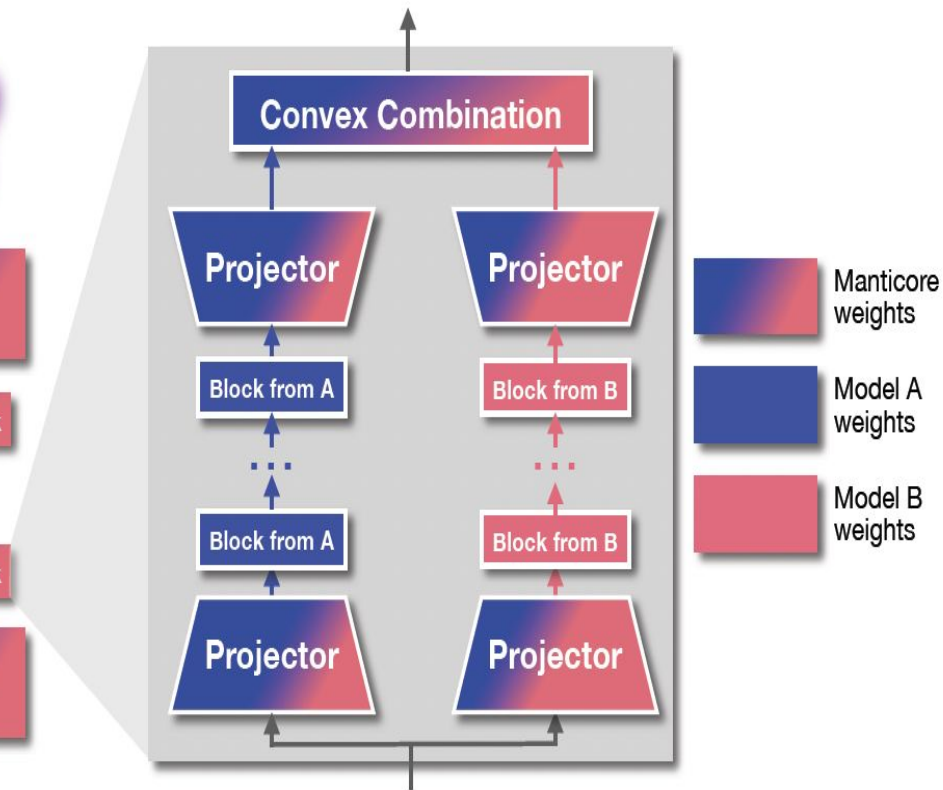
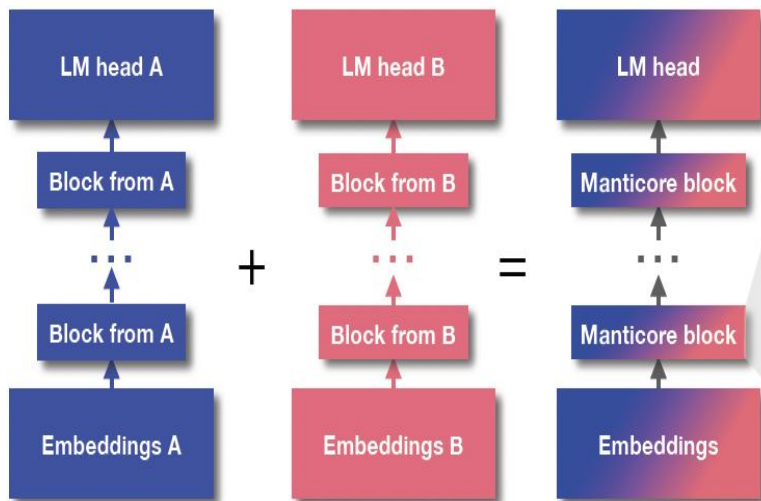
Mamba



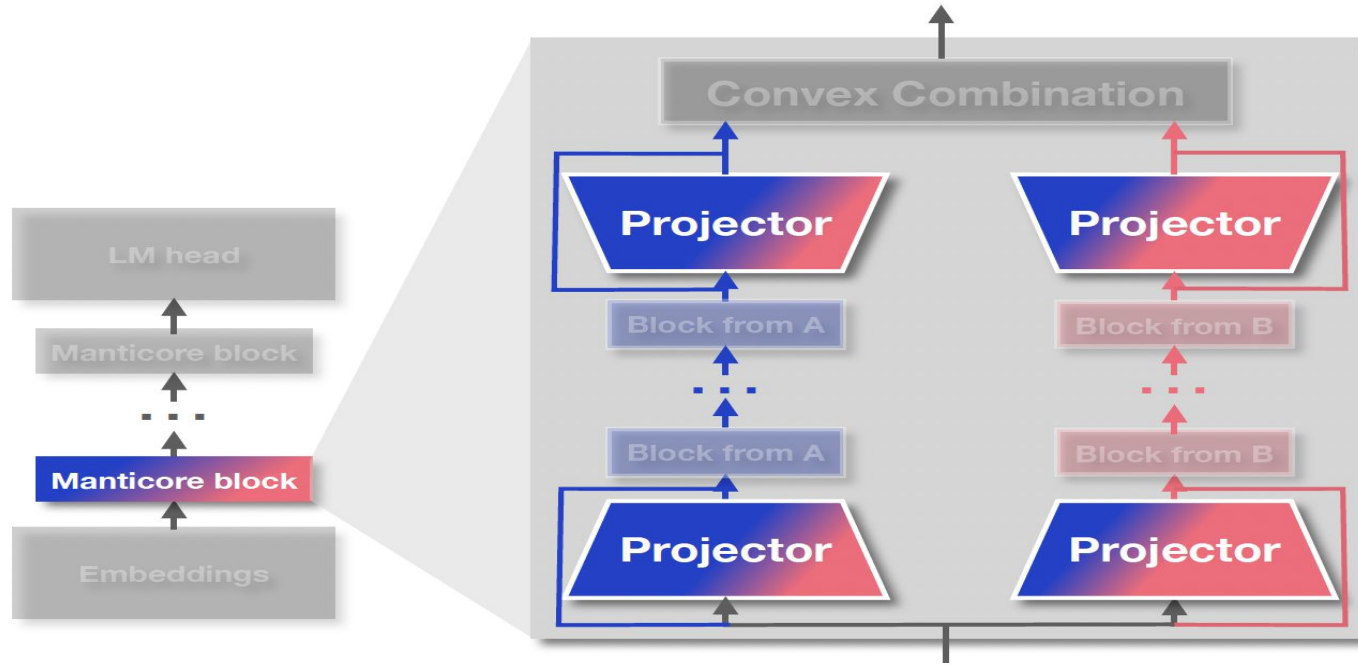
MambaFormer




MANTIGORE



Manticore - Projectors



Manticore - Convex combination



Manticore - Equations

$$\text{Proj-in}_{(k)}^{(\ell)}(x; \alpha) := (1 - \alpha) \cdot \text{Linear}_{d_{\max} \rightarrow d_{M_{(k)}}}(x) + \alpha \cdot \text{Trunc}(x; d_{M_{(k)}})$$

$$\text{Proj-out}_{(k)}^{(\ell)}(x; \alpha) := (1 - \alpha) \cdot \text{Linear}_{d_{M_{(k)}} \rightarrow d_{\max}}(x) + \alpha \cdot \text{Pad}(x; d_{\max}).$$

$$h_k(x; \alpha_k, I_k) = \left(\text{Proj-out}_{(k)}^{(I_k, n_k)} \circ M_{(k)\text{Block}}^{(I_k, n_k)} \circ \dots \circ M_{(k)\text{Block}}^{(I_k, 1)} \circ \text{Proj-in}_{(k)}^{(I_k, 1)} \right) (x; \alpha_k).$$

$$\text{Mix}_{\alpha}(x; I_1, \dots, I_K) = \sum_{k \in [K]} \alpha_k h_k(x; \alpha_k, I_k).$$

Implementation Details

Created our own codebase from scratch with the following modular components:

- Gated Combiners (with convex combination)
- Gated Splitters
- Hybrid Model class (with task-specific output heads)
- Trainer

Next, trying to utilize DeepSpeed's ZeRO optimization or other model parallelism approaches to work with larger model backbones

IMDB Results

With 3 epochs, 2500 iterations/ epoch, we got the following dev accuracies:

Transformer	Mamba	Hybrid
58.7	89.5	84.8 (incomplete training)

MAD Tasks

- Selective Copying

input example: a c [b] t [b] [i] [i] [b] [i] | a c [b] t [b] a c [b] t

- In-context Recall

input example: a b d e f g | a b f g

- Noisy In-context Recall

input example: a b h d e f g | i a b f g

- Fuzzy In-context Recall

input example: (a d) (b) (d a f) (e g) | (d a f) (e g)

- Memorization

key-value dictionary example: {a:b, c:d, e:f} input example: a [i] c [i] e [i] a [i] | a b c d e f a b

MAD Task Config

- **MAD config**

- Number of layers: 2
- Vocab size: 16
- Seq len: 32
- Train examples: 4096
- Test examples: 256
- Hidden size: 64
- Number of Transformer Heads: 4
- Noise fraction: 0.2

- **Training**

- Epochs: 20
- Batch size: 32
- Learning rate: $5e-4$

MAD Tasks Results

Table: Best Test Accuracy on MAD Tasks

Task	Transformer	Mamba	Hybrid 1-blk	Hybrid 2-blk	MambaFormer
Selective Copying	0.91	1.0	0.99	0.99	0.98
In-context Recall	0.35	1.0	1.0	1.0	0.99
Noisy In-context Recall	0.42	1.0	1.0	0.99	1.0
Fuzzy In-context Recall	0.35	0.36	0.42	0.43	0.15
Memorization	1.0	1.0	1.0	1.0	1.0

References

- Albert Gu, Karan Goel, and Christopher Ré. 2021. Efficiently modeling long sequences with structured state spaces. CoRR, abs/2111.00396.
- Albert Gu and Tri Dao. 2024. Mamba: Linear-time sequence modeling with selective state spaces.
- Jongho Park, et al. 2024. Can mamba learn how to learn? a comparative study on in-context learning tasks.
- Michael Poli, et. al. 2024. Mechanistic design and scaling of hybrid architectures.
- MohammadReza Davari and Eugene Belilovsky. 2024. Model breadcrumbs: Scaling multi-task model merging with sparse masks.
- Nicholas Roberts, et. al. 2024. Pretrained hybrids with mad skills.
- Takuya Akiba, et. al. Evolutionary optimization of model merging recipes.